FEAT-ZD-48600 · DES-04 Task Creation · Engineering tasks with goal, AC, dependencies and size budget

Design agent · item feat.tasks · source hobasa-design-agent/outputs/FEAT-ZD-48600/feat.tasks.md · approved by Developer sim.developer

feat.tasks: FEAT-ZD-48600

Engineering tasks for FEAT-ZD-48600 (ZD-48600, "Show which overdue invoices belong to vendors that usually pay late"). This is DES-04 Task Creation (attempt 1), a draft for the Developer gate.

Each user story in feat.stories gets one engineering task. Each task has a goal, the components it touches, the acceptance criteria it must meet, the tests to add, its dependencies and a size budget. Tasks follow the approved design (feat.design Part B and Part C, DT-1 to DT-7) and the developer answers in feat.clarifications. They add no scope beyond feat.stories STORY-01 to STORY-08 and choose no new technology: only the Python 3.9 standard library and Node.js 22 CommonJS (feat.design DD-1). Statements marked **agent judgement** are this agent's proposal; no input states them as fact.

Task IDs (TSK-nn) are labels assigned in this step. They are not taken from any input.

Field	Value	Source
Trace ID	FEAT-ZD-48600	record:brief
Step	DES-04 Task Creation (attempt 1)	record:brief
Trigger	Design document with HLD & LLD	record:brief
Owner / gate	Architect / Developer	record:brief
Inputs	feat.stories (0590ea4), feat.design (722a98f), feat.clarifications (36d4dac), org.risk_map (a6b3e2a)	record:brief
Companion output	feat.tier (T2)	record:brief; org.risk_map
Code read	The four repositories' source and tests from the step's tool view	tool:repo
Coding plan	DE + DS + BE + FE in parallel, then the Testing Agent end-to-end proof	feat.design (Part C, last line); feat.clarifications (CL-X1, CL-T5)

1. Task list

Task	Story	Lane / repository	Goal	Components	Size	Depends on	Ready
TSK-01	STORY-01	Data Engineering / ap_pipeline	Accept the optional `paid_date` and reject invalid values with reason `'paid_date'`	HC-1 `transform.to_model`, `_parse_paid_date`, `PAID_DATE_REASON`	M	— (feed v3 gates release only)	Yes
TSK-02	STORY-02	Data Engineering / ap_pipeline	Build each vendor's payment-delay history	HC-2 `delays.delays_by_vendor` (new module)	S	TSK-01	Yes, after TSK-01
TSK-03	STORY-03	Data Science / ap_forecast	Rule-based vendor risk level, no LLM	HC-4 `risk.vendor_risk` (new module), `__all__`	S	—	Yes

Task	Story	Lane / repository	Goal	Components	Size	Depends on	Ready
TSK-04	STORY-04	Data Science / ap_forecast	State the risk level in the explanation prompt when given	HC-5 `explain.build_explanation_request(..., risk=None)`	S	— (PROV-ACME-2026-11 order gates release only)	Yes
TSK-05	STORY-05	Backend / ap_aging	List unpaid invoices more than 30 days past due with vendor risk	HC-6 `aging.overdue_with_risk`, `OVERDUE_AFTER_DAYS`, `TWO_PLACES`, `__all__`	M	—	Yes
TSK-06	STORY-06	Backend / ap_aging	Serialise the overdue rows to the CT-1 JSON contract	HC-7 `aging.overdue_rows_to_json`, `__all__`	S	TSK-05	Yes, after TSK-05
TSK-07	STORY-07	Frontend / ap_portal	Render the overdue risk table SC-1	HC-8 `renderOverdueRisk`, `escapeHtml`, `toTwoDecimalString`, `riskClass`, `RISK_CLASSES`, `OVERDUE_HEADINGS`	M	CT-1 field list only (not TSK-06 code)	Yes
TSK-08	STORY-08	Testing Agent / cross-repository	Prove feed → delays → risk → overdue list → portal table	HC-10 end-to-end harness (test asset, not deployed)	M	TSK-01, 02, 03, 05, 06, 07 merged	After the four lanes merge

Sources: feat.stories (sections 1, 2, 4); feat.design (A2, Part C, B5); feat.clarifications (section 10). Sizes S and M for TSK-01 to TSK-07 are carried from feat.design Part C ("agent judgement from the number of functions and tests per lane"). The size M for TSK-08 is this step's agent judgement: the harness spans two runtimes and four packages.

```
flowchart LR
  T01[TSK-01 paid_date] --> T02[TSK-02 delays_by_vendor]
  T05[TSK-05 overdue_with_risk] --> T06[TSK-06 CT-1 JSON]
  T06 -. CT-1 contract .-> T07[TSK-07 renderOverdueRisk]
  T03[TSK-03 vendor_risk]
  T04[TSK-04 explanation risk]
  T02 --> T08[TSK-08 end to end]
  T03 --> T08
  T06 --> T08
  T07 --> T08
  X1([feed v3: PQ-6]) -. release .-> T01
  X2([PROV-ACME-2026-11: PQ-5]) -. release .-> T04
```

2. Size budget

The size classes are relative, as in feat.design Part C. This step turns each class into a scope budget: the files a task may change and what it must leave alone. A task that needs a file outside its budget, or a change to an existing signature, stops and goes back to the Architect (agent judgement; the signature rule is org.architecture via feat.design: "Keep every existing signature; add new functions.").

Class	Budget (agent judgement)
S	One new function, or one additive parameter, in one source module, plus its unit tests, any `__init__` export and, where section 2 assigns it, the package version step

Class	Budget (agent judgement)
M	A behaviour change to an existing function or a new function with private helpers and error branches, in one source module, plus its unit tests, exports and the package version step

No numeric estimate (hours or points) is set. Backlog estimates are the PM's (feat.stories SQ-1), and no input gives a numeric budget. This is open question TKQ-1.

CI bar for every coding task (TSK-01 to TSK-07). Compile, unit tests, coverage of at least 90% and static checks stay green. Python: 3.9 syntax only (no `match`, no `X | Y` unions), no `print`, `eval`, `exec` or bare `except`, lines at most 140 characters. ap_portal: no line matching `console.log` or `eval(` (comments included), lines at most 140 characters. Every error branch needs its own test, because a new module with no tests lowers the package's coverage figure (feat.design NF-9; feat.clarifications CL-X3, CL-X5, CL-X6).

Branch and commit. Branch `feature/<story>` (the repositories also have `main`, released, and `dev`, integration); commit messages start with the story ID (for example `STORY-01`), and the task ID may follow it (feat.clarifications CL-X7; record:memory conventions). This answers feat.clarifications CQ-5 for this step: the story label stays the commit prefix, and the task ID is added for traceability (agent judgement).

Version step (one per package). Each changed package takes one minor version step, for the lane to confirm (feat.design DD-11; feat.clarifications CL-X8). To avoid two tasks bumping the same package, this step assigns each bump to one task (agent judgement). Current versions, which the owning task steps up: ap_pipeline `0.4.0` today → TSK-01; ap_forecast `0.9.0` today → TSK-03; ap_aging `1.4.0` today → TSK-05; ap_portal `2.3.0` today (in `src/ap_portal/index.js` and `package.json`) → TSK-07. The next version number is for the lane to confirm.

Test files (to limit merge conflicts between parallel tasks). Agent judgement within feat.design B1.5 to B4.5 and feat.clarifications CL-X4, which allow either the existing test file or a new `tests/test*.py` file: TSK-01 → `tests/test_pipeline.py`; TSK-02 → new `tests/test_delays.py`; TSK-03 → new `tests/test_risk.py`; TSK-04 → `tests/test_forecast.py`; TSK-05 and TSK-06 → new `tests/test_overdue.py`; TSK-07 → `tests/portal.test.js`.

3. Tasks

TSK-01 Accept the optional paid_date in the vendor feed (STORY-01)

Field	Value	Source
Story / lane	STORY-01 / Data Engineering, ap_pipeline	feat.stories (STORY-01)
Goal	`to_model` reads an optional `paid_date`, adds a `paid_date` key (a `date` or `None`) to every record, and rejects a non-blank invalid value with reason `'paid_date'`; feed v2 still loads	feat.stories (STORY-01); feat.design (HC-1, B1.2)
Files in budget	`src/ap_pipeline/transform.py`; `src/ap_pipeline/__init__.py` (version step only); `tests/test_pipeline.py`	feat.design (B1.1, B1.5)
Must not change	`ingest.py`, `dq.py`, `pipeline.py`, `schedule.json`; the `to_model` signature and `(records, rejected)` return shape; any existing key or value; `due_date` parsing (no added `strip()`)	feat.design (B1.1, HC-3); feat.stories (AC-01.8); feat.clarifications (CL-DE3)
Size	M	feat.design (Part C, DT-1)
Depends on	None in code. Feed v3 in production (PQ-6) gates release, not coding; build on synthetic v3 rows	feat.clarifications (CL-DE11); feat.design (Part C "External dependencies")

Implementation notes

1. Keep the existing `try` block exactly as it is. Parse `paid_date` only for a row that passed it: `_parse_paid_date(r.get("paid_date"))`. If the parse sat inside the existing `try`, the reason would become `"ValueError"` and AC-01.4 would fail (feat.clarifications CL-DE1; feat.design B1.2 steps 1 to 3, DD-6).
2. Missing column, `None`, blank or whitespace-only → `None`; otherwise `date.fromisoformat(value.strip())` (feat.clarifications CL-DE2).
3. Unpaid row with a valid `paid_date`: accept it and keep the date (interim DD-7, OQ-B6).

Acceptance criteria: AC-01.1 to AC-01.8 (feat.stories STORY-01). Pending markers kept: AC-01.6 (TQ-10), AC-01.7 (OQ-B6) (feat.clarifications section 10).

Tests to add: `test_paid_date_empty_is_none`, `test_feed_v2_without_paid_date_column`, `test_paid_date_parsed_for_paid_row`, `test_invalid_paid_date_rejected` (values `"2026-13-01"` and `"not-a-date"`; no compact date forms, CL-DE4), `test_existing_reason_wins`, `test_unpaid_row_keeps_paid_date`. Existing tests pass unchanged (feat.design B1.5; feat.clarifications CL-DE4, CL-DE10).

Done when: every listed test passes, the CI bar holds, and the ap_pipeline version step is proposed (feat.stories TK-01.4, TK-01.5).

TSK-02 Build each vendor's payment-delay history (STORY-02)

Field	Value	Source
Story / lane	STORY-02 / Data Engineering, ap_pipeline	feat.stories (STORY-02)
Goal	New pure function `delays_by_vendor(records: Iterable[Dict]) -> Dict[str, List[int]]`: `(paid_date - due_date).days` for paid records with a `paid_date`, in feed order, vendors in first-seen order, negative values kept	feat.design (HC-2, B1.2); feat.clarifications (X-1, X-2)
Files in budget	New `src/ap_pipeline/delays.py`; new `tests/test_delays.py`	feat.design (B1.1, B1.5)
Must not change	`transform.py` (beyond TSK-01); `ap_pipeline/__init__.py` exports (no `__all__` added; import as `from ap_pipeline.delays import delays_by_vendor`)	feat.clarifications (CL-DE9); feat.design (DQ-4)
Size	S	feat.design (Part C, DT-2)
Depends on	TSK-01 (reads the `paid_date` key)	feat.stories (section 2); feat.design (Part C "Lane dependencies")

Implementation notes: skip unless `record.get("paid")` is truthy and `record.get("paid_date") is not None`; then `result.setdefault(record["vendor"], []).append(...)`. No I/O; do not change `records` (feat.design B1.2; feat.clarifications X-2).

Acceptance criteria: AC-02.1 to AC-02.8 (feat.stories STORY-02). Pending marker kept: AC-02.6 (OQ-B6).

Tests to add: `test_delays_by_vendor_*`: paid-with-date only; feed order; negative kept; paid-without-date excluded; unpaid-with-date excluded; vendor with no qualifying record has no key; input list unchanged (feat.design B1.5).

Done when: tests pass and the CI bar holds, including coverage of the new module (feat.stories TK-02.3; feat.clarifications CL-X6).

TSK-03 Give each vendor a rule-based risk level (STORY-03)

Field	Value	Source
Story / lane	STORY-03 / Data Science, ap_forecast	feat.stories (STORY-03)
Goal	New `vendor_risk(history: Iterable[int]) -> str`: `'high'` if `predict_delay_days(history) > 30`, `'medium'` if `> 10`, else `'low'`; no LLM, no network, no `models.json` read	feat.design (HC-4, B2.2, DD-9)
Files in budget	New `src/ap_forecast/risk.py` (constants `HIGH_ABOVE_DAYS = 30`, `MEDIUM_ABOVE_DAYS = 10`); `src/ap_forecast/__init__.py` (import, `__all__`, version step); new `tests/test_risk.py`	feat.design (B2.1, B2.2, B2.5)
Must not change	`predict.py`, `models.json`; `risk.py` imports only `predict_delay_days`, not `explain`	feat.design (B2.1, B2.2); feat.clarifications (CL-DS4)
Size	S	feat.design (Part C, DT-3)
Depends on	None (uses the existing `predict_delay_days`)	feat.stories (STORY-03)

Implementation notes: call `predict_delay_days(history)` exactly once, so a generator works (feat.clarifications CL-DS2). Do not validate history values; `predict_delay_days` already applies `int(d)` (CL-DS3).

Acceptance criteria: AC-03.1 to AC-03.8 (feat.stories STORY-03). Pending markers kept: AC-03.5 (PQ-1), AC-03.6 (TQ-5).

Tests to add: `test_vendor_risk_levels`, `test_vendor_risk_boundaries` (single-value histories `[30]`, `[31]`, `[10]`, `[11]`, so the rounded prediction equals the value; record:memory L-007), `test_vendor_risk_empty_and_early`, `test_vendor_risk_no_llm` (patch `ap_forecast.explain.load_config` to fail), `test_vendor_risk_exported` (feat.design B2.5; feat.clarifications CL-DS1, CL-DS4).

Done when: tests pass, the CI bar holds, and the ap_forecast version step is proposed (shared with TSK-04, owned here) (feat.stories TK-03.4).

TSK-04 State the risk level in the payment-delay explanation (STORY-04)

Field	Value	Source
Story / lane	STORY-04 / Data Science, ap_forecast	feat.stories (STORY-04)
Goal	`build_explanation_request(vendor, predicted_days, config=None, risk=None)`: unchanged when `risk is None`; otherwise today's prompt followed by `f" Risk level: {risk}."`	feat.design (HC-5, B2.2)
Files in budget	`src/ap_forecast/explain.py`; `tests/test_forecast.py`	feat.design (B2.1, B2.5)
Must not change	Provider, model and `max_tokens` source (`config or load_config()`); `models.json`; positional use of `config`	feat.stories (AC-04.3, AC-04.4); feat.design (B2.2)
Size	S	feat.design (Part C, DT-4)
Depends on	None in code. Order against PROV-ACME-2026-11 (PQ-5) gates release: whoever merges second rebases and re-runs T-7. Sending the label to `acme-llm` awaits Architect confirmation (TQ-8), which gates release, not coding	feat.clarifications (CL-DS7, CL-DS8); feat.design (SE-3)

Implementation notes: any non-`None` value is appended without validation (interim IF-4.5; an empty string gives `" Risk level: ."`, CQ-4). Nothing new calls this function with `risk`; it is not in the end-to-end wiring (feat.clarifications CL-DS5, CL-DS6).

Acceptance criteria: AC-04.1 to AC-04.5 (feat.stories STORY-04). Pending marker kept: AC-04.5 (TQ-4).

Tests to add: `test_request_with_risk` (prompt ends with `' Risk level: high.'`), `test_request_without_risk_unchanged` (with and without explicit `risk=None`). Existing tests `test_predict`, `test_request_uses_config`, `test_request_with_explicit_config` pass unchanged (feat.design B2.5).

Done when: tests pass and the CI bar holds (feat.stories TK-04.3).

TSK-05 List unpaid invoices more than 30 days past due with vendor risk (STORY-05)

Field	Value	Source
Story / lane	STORY-05 / Backend, ap_aging	feat.stories (STORY-05)
Goal	New pure `overdue_with_risk(invoices, as_of, risk_by_vendor) -> List[Dict]`: unpaid invoices with `days_past_due > 30`, rows with exactly `id`, `vendor`, `days`, `amount` (`Decimal`, 2 dp), `risk` (`'unknown'` when the vendor is not in the map), sorted by `days` descending then `id` ascending	feat.design (HC-6, B3.2)
Files in budget	`src/ap_aging/aging.py` (add `OVERDUE_AFTER_DAYS = 30`, `TWO_PLACES = Decimal("0.01")`, `List` typing, the function); `src/ap_aging/__init__.py` (export, version step); new `tests/test_overdue.py`	feat.design (B3.1, B3.2, B3.5)
Must not change	`aging_buckets`, `days_past_due`, `bucket_for`, `BUCKETS`. Never call `bucket_for` (`bucket_for(30)` returns `"31-60"`)	feat.design (DD-4, B3.1); feat.clarifications (CL-BE1); tool:repo (`aging.py`)
Size	M	feat.design (Part C, DT-5)
Depends on	None in code	feat.stories (section 2)

Implementation notes: skip `inv.get("paid")` truthy (D-2026-050); `as_of` is required, with no default (CL-BE2); `Decimal(inv["amount"]).quantize(TWO_PLACES)` with the default context, never `float` (DD-10; record:memory D-2026-044, L-011); missing `vendor` on an overdue unpaid invoice raises `KeyError` (CL-BE4).

Acceptance criteria: AC-05.1 to AC-05.10 (feat.stories STORY-05). Pending markers kept: AC-05.6 (PQ-2), AC-05.10 (TQ-9).

Tests to add (`AS_OF = date(2026, 10, 25)`, due dates built as `AS_OF - timedelta(days=N)` so `days` is 30, 31, 42 and 75): `test_overdue_boundary_30_and_31` (record:memory L-007), `test_overdue_excludes_paid_and_not_due`, `test_overdue_row_keys_and_sort` (tie on 42: `SYN-INV-002` before `SYN-INV-003`), `test_overdue_unknown_risk`, `test_overdue_amount_quantised` (`Decimal("980.5")` → `Decimal("980.50")`), `test_overdue_missing_vendor_raises`, `test_existing_aging_unchanged` (pins `bucket_for(30) == "31-60"`) (feat.design B3.5; feat.clarifications CL-BE2).

Done when: tests pass, the CI bar holds, and the ap_aging version step is proposed (shared with TSK-06, owned here) (feat.stories TK-05.4).

TSK-06 Hand the overdue list to the portal as JSON, contract CT-1 (STORY-06)

Field	Value	Source
Story / lane	STORY-06 / Backend, ap_aging	feat.stories (STORY-06)
Goal	New `overdue_rows_to_json(rows) -> str`: JSON array of `{id, vendor, days (int), amount (two-decimal string), risk}` in CT-1 key order and the given row order; `'[]'` when empty; no float anywhere	feat.design (HC-7, CT-1, B3.2)
Files in budget	`src/ap_aging/aging.py` (add `import json`, the function); `src/ap_aging/__init__.py` (export); `tests/test_overdue.py`	feat.design (B3.1, B3.2)
Must not change	`overdue_with_risk` behaviour from TSK-05; existing ap_aging functions. No transport layer (a transport would re-trigger DES-01)	feat.design (DD-3, A8); feat.clarifications (CL-BE10)
Size	S	feat.design (Part C, DT-6)
Depends on	TSK-05 (takes its rows)	feat.stories (section 2)

Implementation notes: `str(Decimal(r["amount"]).quantize(TWO_PLACES))` (feat.clarifications X-4). Share the CT-1 field list with the Frontend lane and the Testing Agent (feat.stories TK-06.3).

Acceptance criteria: AC-06.1 to AC-06.4 (feat.stories STORY-06). All four stay [Pending TQ-3] (feat.clarifications section 10).

Tests to add: `test_rows_to_json`: `"1250.00"` for SYN-INV-004; every `amount` a two-decimal string; `days` an integer; `[]` for no rows; key order. Compare `json.loads` values, not raw strings (feat.design B3.5; feat.clarifications CL-BE8).

Done when: tests pass and the CI bar holds (feat.stories TK-06.4).

TSK-07 Render the overdue risk table (STORY-07)

Field	Value	Source
Story / lane	STORY-07 / Frontend, ap_portal	feat.stories (STORY-07)
Goal	Exported `renderOverdueRisk(rows)`: `<table class="overdue-risk">`, a plain header row (Invoice, Vendor, Days overdue, Amount), one `<tr class="risk-…">` per row in the given order, amount cell `class="amount"`; `''` for missing, `null` or `[]`	feat.design (HC-8, B4.2); feat.clarifications (CL-FE1, CL-FE2)
Files in budget	`src/ap_portal/index.js` (constants, private helpers, export, `VERSION` step); `package.json` (version only); `tests/portal.test.js`	feat.design (B4.1); feat.clarifications (X-5)
Must not change	`LABELS`, `formatAmount`, `renderAgingSummary` and their tests; never use `formatAmount`, `Number()` or `parseFloat` for the amount; no re-sorting	feat.design (B4.1, DD-5); feat.clarifications (CL-FE2, CL-FE3, CL-FE7)
Size	M	feat.design (Part C, DT-7)
Depends on	CT-1 field list only; can be built in parallel with TSK-06	feat.stories (section 2); feat.design (Part C)

Implementation notes: amount: string only, pattern `^-?\d+(\.\d{1,2})?$`, padded as a string (`"980.5"` → `980.50`); otherwise `TypeError`. `days`: `Number.isInteger`, else `TypeError`. Non-array `rows` (other than missing / `null`) → `TypeError`. `escapeHtml` on `id` and `vendor`; class from the `RISK_CLASSES` allow-list. No risk text, no Risk column, no row actions (feat.design B4.2, B4.4, DD-8; feat.clarifications CL-FE3 to CL-FE5, CL-FE8).

Acceptance criteria: AC-07.1 to AC-07.10 (feat.stories STORY-07). Pending markers kept: AC-07.2 (TQ-7), AC-07.6 (TQ-11, OQ-B10), AC-07.7 (OQ-B8), AC-07.8 (TQ-6), AC-07.9 (OQ-B9, PQ-P7).

Tests to add: from feat.design B4.5: `renders overdue risk table`, `maps every risk value to its class`, `returns empty string for no rows`, `shows amount strings unchanged`, `pads short amounts as strings`, `rejects non-string amounts`, `escapes id and vendor`, `unexpected risk renders risk-unknown`. From feat.clarifications CL-FE6, for coverage: non-array `rows` → `TypeError`; non-integer `days` → `TypeError`; `null` → `''`; amount with an exponent, a thousands separator or three decimals → `TypeError`. Existing tests pass unchanged. Run with `node:test` and `NODE_PATH=src` (CL-FE10).

Done when: tests pass, the CI bar holds, and `VERSION` in `index.js` and `"version"` in `package.json` move together (feat.clarifications X-5; feat.stories TK-07.5).

TSK-08 Prove feed → delays → risk → overdue list → portal table (STORY-08)

Field	Value	Source
Story / lane	STORY-08 / Testing Agent, cross-repository	feat.stories (STORY-08)
Goal	End-to-end harness: `pipeline.run` → `delays_by_vendor` → `vendor_risk` per vendor → `overdue_with_risk` → `overdue_rows_to_json` → Node.js 22 `renderOverdueRisk(JSON.parse(payload))`; the rendered rows match `overdue_with_risk` from the same feed	feat.design (B5); feat.clarifications (CL-T1)
Files in budget	Test assets only; the harness is not a deployed component (making it one would re-trigger DES-01). Location is the Testing Agent's choice (TKQ-3)	feat.design (HC-10, A8)
Must not change	Any product source in the four repositories	feat.design (B5)
Size	M (agent judgement)	section 1
Depends on	TSK-01, TSK-02, TSK-03, TSK-05, TSK-06, TSK-07 merged; TSK-04 is not in the flow	feat.stories (section 2, SA-5); feat.clarifications (CL-T5)

Implementation notes: synthetic feed only (`SYN-INV-*`, "Synthetic Vendor <letter>") with an exactly-30-days invoice, a paid invoice, an invalid `paid_date` row and a vendor with no paid history; a fixed `as_of`, not today (feat.clarifications CL-T2, CL-T3, CL-X9).

Acceptance criteria: AC-08.1 to AC-08.4 (feat.stories STORY-08). Pending marker kept: AC-08.1 (PQ-2).

Done when: AC-08.1 to AC-08.4 are asserted and pass (feat.stories TK-08.3).

4. Traceability

Task	Story	Design task	HLD component	Acceptance criteria	TRD tests
TSK-01	STORY-01	DT-1	HC-1	AC-01.1 to AC-01.8	T-1, T-2, T-3
TSK-02	STORY-02	DT-2	HC-2	AC-02.1 to AC-02.8	T-4
TSK-03	STORY-03	DT-3	HC-4	AC-03.1 to AC-03.8	T-5, T-6
TSK-04	STORY-04	DT-4	HC-5	AC-04.1 to AC-04.5	T-7
TSK-05	STORY-05	DT-5	HC-6	AC-05.1 to AC-05.10	T-8, T-9, T-10 (row part), T-11
TSK-06	STORY-06	DT-6	HC-7	AC-06.1 to AC-06.4	T-10 (JSON part)
TSK-07	STORY-07	DT-7	HC-8	AC-07.1 to AC-07.10	T-12 to T-16
TSK-08	STORY-08	— (B5)	HC-10	AC-08.1 to AC-08.4	T-17

Sources: feat.stories (sections 4, 7); feat.design (A2, Part C, D1). Every story STORY-01 to STORY-08 has exactly one task.

5. Story commitment status

All eight stories are ready to build on the approved design (feat.clarifications section 10). Commitment to an iteration, priority and estimate are not set here: backlog prioritisation is the PM's (feat.stories SQ-1). This file proposes the stories for commitment with their tasks; the Developer gate confirms the task breakdown and sizes.

6. Assumptions (new in this step; agent judgement, please confirm at the gate)

ID	Assumption
TA-1	One engineering task per story, matching feat.design DT-1 to DT-7, plus TSK-08 for the end-to-end story
TA-2	The size budget is a scope budget (files in budget, must not change), not a time estimate
TA-3	Each package's version step is owned by one task (TSK-01, TSK-03, TSK-05, TSK-07)
TA-4	Test file split as in section 2, to reduce merge conflicts between parallel tasks
TA-5	Commit messages keep the story ID as the prefix; the task ID may follow (answers CQ-5 for this step)

7. Open questions (all non-blocking)

ID	Question	Interim answer	Owner (proposed)
TKQ-1	Should tasks carry a numeric estimate (hours or points)? No input gives one	Size class and scope budget only	Developer / PM
TKQ-2	Confirm the version-step owners and test-file split (TA-3, TA-4)	As in section 2	Developer
TKQ-3	Where should the end-to-end harness live? No input names a location	Testing Agent's choice; test asset, not deployed	Testing Agent

Carried, unchanged and non-blocking: TQ-1 to TQ-11 (feat.trd via feat.stories and feat.design), DQ-1 to DQ-5 (feat.design D3), CQ-1 to CQ-4 and CQ-6 (feat.clarifications 9.1), SQ-1 to SQ-4 (feat.stories 6.1), OQ-B1 to OQ-B10, PQ-1 to PQ-6, PQ-P1 to PQ-P7 (feat.stories 6.2). CQ-5 is answered for this step by TA-5.

8. Validation check

Check	Result	Source
Every story has a task with goal, components, acceptance criteria, tests, dependencies and size	Pass (sections 1, 3)	record:brief
No product code changed; no new technology	Pass (DD-1 carried)	feat.design (DD-1)
Every task cites the design, a contract or a story	Pass	feat.design; feat.stories
Decisions D-2026-031, D-2026-044, D-2026-050; lessons L-007, L-011	Consistent (TSK-03, TSK-05, TSK-06, TSK-07)	record:memory
Conflict with an approved output	None. Form differences resolved as in feat.clarifications section 1	feat.clarifications (X-1 to X-6)
Blocking questions	None	—